

## Aws\_Storage\_Service:

Concept: 1 "Amazon Elastic Block Store"

Amazon Elastic Block Store (Amazon EBS) is a cloud-based block storage service offered by Amazon Web Services (AWS). It provides persistent, high-performance storage volumes that can be attached to Amazon EC2 (Elastic Compute Cloud) instances. These volumes function similarly to traditional hard drives, allowing you to store files, run databases, and install applications.

### Key Features of Amazon EBS:

- **Persistent Storage:** Data stored in EBS volumes remains intact even if the associated EC2 instance is stopped or terminated.
- **High Availability:** EBS volumes are automatically replicated within their Availability Zone to protect against hardware failures.
- **Scalability:** You can dynamically scale EBS volumes in terms of size and performance to meet your application needs without downtime.
- **Snapshots:** EBS supports point-in-time snapshots, which are incremental backups stored in Amazon S3. These snapshots can be used for data recovery, migration across regions, or creating new volumes.
- **Encryption:** EBS offers encryption at rest and in transit, managed by AWS Key Management Service.

### EBS Volume Types:

Amazon EBS provides various volume types to cater to different performance and cost requirements:

- **General Purpose SSD (gp2/gp3):** Balanced performance for a wide range of workloads.
- **Provisioned IOPS SSD (io1/io2):** Designed for I/O-intensive applications like large databases.
- **Throughput Optimized HDD (st1):** Ideal for frequently accessed, throughput-intensive workloads.
- **Cold HDD (sc1):** Suitable for infrequently accessed data with lower cost requirements.

### Common Use Cases:

**Databases:** Hosting relational and NoSQL databases that require consistent and low-latency storage

**Enterprise Applications:** Running applications like SAP, Oracle, and Microsoft products that demand high availability and performance

**Backup and Recovery:** Creating snapshots for disaster recovery and data migration purposes

**Big Data Analytics:** Storing and processing large datasets for analytics workloads.

### Ebs\_setup :

Step:1

We have to launch an instance

**Instances (1)**
[Info](#)

Connect

Instance state

Actions

Launch instances

All states

< 1 >

<input type="checkbox"/>	Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone	Public IP
<input type="checkbox"/>	my-server	i-05fed8b4a5a789433	Running	t2.micro	Initializing	<a href="#">View alarms +</a>	us-east-1c	ec2-35-1...

Step:2

Go to Elastic Block Storage Inside we have volumes and create a new

## ▼ Elastic Block Store

Volumes

Snapshots

Lifecycle Manager

volume

Step:3

We have default volume and we have to create a customized volume like hard disk because if we stop an instance volume will delete to overcome these problem we are create an

volume

**Volumes (1)**
[Info](#)

Last updated  
less than a minute ago

Actions

Create volume

Saved filter sets  
Choose filter set

< 1 >

<input type="checkbox"/>	Name	Volume ID	Type	Size	IOPS	Throughput	Snapshot ID	Created
<input type="checkbox"/>		vol-0029193aa56591b21	gp2	8 GiB	100	-	snap-0898870...	2025/07/08 11:49 GMT+5:...

Step:4

Create a volume

## Create volume [Info](#)

Create an Amazon EBS volume to attach to any EC2 instance in the same Availability Zone.

### Volume settings

Volume type [Info](#)

General Purpose SSD (gp3)

Size (GiB) [Info](#)

2

Min: 1 GiB, Max: 16384 GiB.

## IOPS | Info

3000

Min: 3000 IOPS, Max: 16000 IOPS.

Input output per second : searching speed

## Throughput (MiB/s) | Info

125

Min: 125 MiB, Max: 1000 MiB. Baseline: 125 MiB/s.

Cancel

Create volume

Step:5

We have to select and go to action and click attach volume

The screenshot shows the AWS Management Console 'Volumes' page. At the top, there's a 'Volumes (1/2)' header with an 'Info' link. Below it is a 'Saved filter sets' section with a 'Choose filter set' dropdown and a search bar. A table lists two volumes:

	Name	Volume ID	Type	Size	IOPS	Throughput
<input type="checkbox"/>		vol-0029193aa56591b21	gp2	8 GiB	100	-
<input checked="" type="checkbox"/>		vol-0d92a35ffbf6b6d15	gp3	2 GiB	3000	125


Below the table, the 'Volume ID: vol-0d92a35ffbf6b6d15' is displayed. There are tabs for 'Details', 'Status checks', 'Monitoring', and 'Tags'. On the right, an 'Actions' menu is open, showing options: 'Modify volume', 'Create snapshot', 'Create snapshot lifecycle policy', 'Delete volume', 'Attach volume', 'Detach volume', 'Force detach volume', 'Manage auto-enabled I/O', 'Manage tags', and 'Fault injection'. The 'Attach volume' option is highlighted. At the top right, there's a 'Create volume' button and a 'Last updated less than a minute ago' status.

## Attach volume | Info

Attach a volume to an instance to use it as you would a regular physical hard disk drive.

### Basic details

#### Volume ID

 vol-0d92a35ffbf6b6d15

#### Availability Zone

us-east-1c

## Instance | Info

Search instance ID or name tag

Only instances in the same Availability Zone as the selected volume are displayed.

Device name [Info](#)

/dev/sdf

Recommended device names for Linux: /dev/xvda for root volume, /dev/sd[f-p] for data volumes.

**i** Newer Linux kernels may rename your devices to **/dev/xvdf** through **/dev/xvdp** internally, even when the device name entered here (and shown in the details) is **/dev/sdf** through **/dev/sdp**.

## Click Attach Volume

Successfully attached volume vol-0d92a35ffb6b6d15 to instance i-05fed8b4a5a789433.

**Volumes (2)** [Info](#) Last updated 1 minute ago [Actions](#) [Create volume](#)

Saved filter sets [Choose filter set](#)

<input type="checkbox"/>	Name <a href="#">↗</a>	Volume ID	Type	Size	IOPS	Throughput	Snapshot ID	Created
<input type="checkbox"/>		vol-0029193aa56591b21	gp2	8 GiB	100	-	snap-0898870...	2025/07/08 11:49 GMT+5:...
<input type="checkbox"/>		vol-0d92a35ffb6b6d15	gp3	2 GiB	3000	125	-	2025/07/08 12:18 GMT+5:...

In use

We have to check data will store in 2gb or not  
Go to server and check

- Connect to instance : `sudo su`
- `[root@ip-172-31-80-53 ec2-user]# df -h`  
Filesystem    Size    Used Avail Use% Mounted on  
devtmpfs     468M    0 468M   0% /dev  
tmpfs        477M    0 477M   0% /dev/shm  
tmpfs        477M   408K 476M   1% /run  
tmpfs        477M    0 477M   0% /sys/fs/cgroup  
**/dev/xvda1    8.0G 1.9G 6.2G 24% /**  
tmpfs        96M    0 96M   0% /run/user/1000
- `[root@ip-172-31-80-53 ec2-user]# lsblk` : list of block device  
NAME    MAJ:MIN RM SIZE RO TYPE MOUNTPOINT  
xvda    202:0    0   8G 0 disk  
└─xvda1 202:1    0   8G 0 part /  
**xvdf    202:80   0   2G 0 disk**
- To check xvdf have file system or not  
`[root@ip-172-31-80-53 ec2-user]# file -s /dev/xvdf`  
/dev/xvdf: data
- We should create a file system  
`[root@ip-172-31-80-53 ec2-user]# mkfs -t xfs /dev/xvdf`  
meta-data=/dev/xvdf            isize=512    agcount=4, agsize=131072 blks  
=                    sectsz=512    attr=2, projid32bit=1  
=                    crc=1       finobt=1, sparse=1, rmapbt=0  
=                    reflink=1    bigtime=0 inobtcount=0  
data    =                    bsize=4096    blocks=524288, imaxpct=25  
=                    sunit=0     swidth=0 blks  
naming    =version 2            bsize=4096    ascii-ci=0, ftype=1  
log       =internal log       bsize=4096    blocks=2560, version=2  
=                    sectsz=512    sunit=0 blks, lazy-count=1  
realtime =none            extsz=4096    blocks=0, rtextents=0
- To check xvdf have file system or not  
`[root@ip-172-31-80-53 ec2-user]# file -s /dev/xvdf`  
**/dev/xvdf: SGI XFS filesystem data (blksz 4096, inosz 512, v2 dirs)**

- We create a folder because we should mount the file system to create folder  
mkdir /ebsdemo

mount /dev/xvdf /ebsdemo

[root@ip-172-31-80-53 ec2-user]# df -h

Filesystem	Size	Used	Avail	Use%	Mounted on
devtmpfs	468M	0	468M	0%	/dev
tmpfs	477M	0	477M	0%	/dev/shm
tmpfs	477M	464K	476M	1%	/run
tmpfs	477M	0	477M	0%	/sys/fs/cgroup
/dev/xvda1	8.0G	1.9G	6.2G	24%	/
tmpfs	96M	0	96M	0%	/run/user/1000
tmpfs	96M	0	96M	0%	/run/user/0
/dev/xvdf	2.0G	47M	2.0G	3%	/ebsdemo

- We have to set-up these settings otherwise data will store in default volume

[root@ip-172-31-80-53 /]# blkid

/dev/xvda1: LABEL="/" UUID="6e43975a-1916-4f3f-b40f-cfed25542d46" TYPE="xfs"  
PARTLABEL="Linux" PARTUUID="8642b54c-0495-4673-bd96-dfabd01ffa3a"

/dev/xvdf: UUID="6fa68314-91a7-46cb-89fc-d3d4d39b9363" TYPE="xfs"

vim /etc/fstab

#

UUID=6e43975a-1916-4f3f-b40f-cfed25542d46 / xfs defaults,noatime 1 1

UUID=6fa68314-91a7-46cb-89fc-d3d4d39b9363 /ebsdemo xfs defaults,nofail 0 0

mount -a to check whether it mount or not

Mainly stop instance and detach volume and after we create a snapshot and we can create a number of volumes and we can change az also

This snapshot is for volume

Create a snapshot ----- go to snapshot and select and go to action and create a volume

- [ec2-user@ip-172-31-85-90 ~]\$ sudo su

- [root@ip-172-31-85-90 ec2-user]# df -h

Filesystem	Size	Used	Avail	Use%	Mounted on
devtmpfs	468M	0	468M	0%	/dev
tmpfs	477M	0	477M	0%	/dev/shm
tmpfs	477M	400K	476M	1%	/run
tmpfs	477M	0	477M	0%	/sys/fs/cgroup
/dev/xvda1	8.0G	1.9G	6.2G	24%	/
tmpfs	96M	0	96M	0%	/run/user/1000

- [root@ip-172-31-85-90 ec2-user]# df -h

Filesystem	Size	Used	Avail	Use%	Mounted on
devtmpfs	468M	0	468M	0%	/dev
tmpfs	477M	0	477M	0%	/dev/shm
tmpfs	477M	404K	476M	1%	/run
tmpfs	477M	0	477M	0%	/sys/fs/cgroup
/dev/xvda1	8.0G	1.9G	6.2G	24%	/
tmpfs	96M	0	96M	0%	/run/user/1000

- [root@ip-172-31-85-90 ec2-user]# lsblk

NAME MAJ:MIN RM SIZE RO TYPE MOUNTPOINT

```

xvda 202:0 0 8G 0 disk
└─xvda1 202:1 0 8G 0 part /
xvdf 202:80 0 2G 0 disk
  • [root@ip-172-31-85-90 ec2-user]# file -s /dev/xvdf
/dev/xvdf: SGI XFS filesystem data (blksz 4096, inosz 512, v2 dirs)
[root@ip-172-31-85-90 ec2-user]# mkdir /ebsdemo
[root@ip-172-31-85-90 ec2-user]# mount /dev/xvdf /ebsdemo
[root@ip-172-31-85-90 ec2-user]# cd /ebsdemo/
  • [root@ip-172-31-85-90 ebsdemo]# ll
total 4
-rw-r--r-- 1 root root 22 Jul 8 07:25 file1
-rw-r--r-- 1 root root 0 Jul 8 07:24 file2
-rw-r--r-- 1 root root 0 Jul 8 07:24 file3
-rw-r--r-- 1 root root 0 Jul 8 07:24 file4
-rw-r--r-- 1 root root 0 Jul 8 07:24 file5
  • [root@ip-172-31-85-90 ebsdemo]# cat file1
this is done by file1

```

Region wise means we take snapshot copy

## What Is an Instance Store?

Instance stores provide temporary storage that is physically attached to the host machine. Unlike EBS volumes—which reside on separate machines and are accessed via protocols like iSCSI—instance stores enable your EC2 instance to access a local physical drive directly on the host. Their temporary nature means that the stored data is lost if the instance is moved to a different host.

## How Instance Stores Work

Consider a scenario where multiple EC2 instances run on the same host machine. When an EC2 instance utilizes an instance store:

- The instance directly accesses a physical drive on that host.
- Provided the instance remains on the same host, even after a reboot, it retains access to its instance store data.

However, when an instance is shut down and later restarted on a different host, it loses access to the original instance store and all associated data because the new host has a different storage configuration.

### Warning

Moving an EC2 **instance** to a different host results in the loss of any data stored on the original **instance** store. Use **instance** stores only for data that can be safely discarded.

## Use Cases and Limitations

Instance stores are best suited for applications where temporary data storage is acceptable and data persistence is not required. Common use cases include:

- Caching
- Buffering
- Temporary file storage

When designing systems, keep in mind that instance stores are not reliable for storing persistent data, as data loss will occur if the instance is migrated.

To summarize:

- Instance stores provide block-level, temporary storage for EC2 instances.
- Data on an instance store is lost if the instance migrates to a different host.
- They are ideal for transient data such as caches or scratch files, but not for persistent storage.

## Launching an Instance with an Instance Store

To begin, navigate to the EC2 console and launch a new instance. For this demonstration, name your instance "instance store demo" and select the Amazon Linux 64-bit AMI. Note that not all EC2 instance types support instance stores; for example, t2.micro and other free tier instances lack instance store volumes. Select an instance type that includes instance store support—even if it comes at a small cost for this short demo.

After choosing an appropriate instance type, select any key pair of your preference.



Review the configuration details. You will notice that the default root volume is 8 GB, and further down, you will see an instance store volume (approximately 75 GB, reported as 69.8 GB). The device name for the instance store is typically something like `/dev/nvme0n1` or `/dev/nvme1n1`. Also, ensure that "Auto-assign Public IP" is enabled so you can easily connect to your instance.

Once your instance is deployed, return to the EC2 console's instance tab. You should now see the instance with its assigned public IP.

### Connecting and Configuring the Instance Store

After noting the public IP, SSH into your instance. Once connected, run the following command to verify the block devices:

```
lsblk
```

You'll see output similar to:

```
NAME          MAJ:MIN RM  SIZE RO TYPE MOUNTPOINTS
nvme1n1       259:0    0 69.8G  0 disk
nvme0n1       259:1    0   8G   0 disk
├─nvme0n1p1   259:2    0   8G   0 part /
└─nvme0n1p128 259:4    0  10M   0 part
```

Here, `nvme0n1` is your root volume (with partition `nvme0n1p1` mounted as `/`), and `nvme1n1` is the instance store volume.

Verify if a file system exists on the instance store volume by running:

```
sudo file -s /dev/nvme1n1
```

If the output shows "data," no file system exists. Create an XFS file system on the volume as follows:

```
sudo mkfs -t xfs /dev/nvme1n1
```

Then, validate the file system creation by running:

```
sudo file -s /dev/nvme1n1
```

You should now see that an SGI XFS filesystem is present on the device.

Next, create a directory to serve as the mount point and mount the instance store with these commands:

```
sudo mkdir /instance-demo  
sudo mount /dev/nvme1n1 /instance-demo/
```

Confirm that the volume is mounted by executing:

```
df -h
```

Now, navigate into the mount directory and create a test file to verify that data is being written to the instance store:

```
cd /instance-demo  
echo "test" | sudo tee test
```

The file "test" is now stored on your instance store volume.

#### Tip

Keep in mind that the instance store is optimized for fast, temporary storage. It is ideal for cache, buffers, or temporary files, but not for data that requires

```
persistence.
```

## Demonstrating the Ephemeral Nature of Instance Store Data

Data on an instance store is ephemeral. A simple reboot will not change the physical host, so your instance store volume and its data remain intact. You can verify this by rebooting the instance from the EC2 console (right-click the instance and select "Reboot"). After the reboot, the public IP address will remain the same.

However, when you stop and then start your instance, it is relocated to a different physical host. This process replaces the instance store with a new, empty volume. To observe this behavior:

1. Record the current public IP address.
2. Stop the instance using the EC2 console.
3. Wait a few minutes, then start the instance again.

After the instance restarts, note the change in the public IP address, which confirms it now runs on a different physical host. SSH into the instance again and run:

```
lsblk
```

You will notice the instance store volume (nvme1n1) is still present. However, if you check the file system with:

```
sudo file -s /dev/nvme1n1
```

and list the contents of /instance-demo, you will find that the test file created earlier is missing. This confirms that data from the previous instance store has been lost due to the instance migration.

### Warning

Do not use instance stores for storing critical data. Always rely on persistent storage solutions like Amazon EBS for data that must be retained.

## Quick Reference Table

Command/Action	Description
lsblk	List block devices
sudo file -s /dev/nvme1n1	Check for a file system on the instance store volume
sudo mkfs -t xfs /dev/nvme1n1	Create an XFS file system on the instance store
sudo mkdir /instance-demo	Create a mount directory
sudo mount /dev/nvme1n1 /instance-demo/	Mount the instance store

```
df -k
```

Confirm the mount status

```
sudo tee test`
```

```
`echo "test"
```